

Random Topics

Last time...

Logistic Regression

Today

- Review & finish logistic regression
- Weighted Least Squares
- Work on HW 5

Maximum Likelihood Estimation

- Logistic regression is used when our outcomes are categorical
- In logistic regression we are not so lucky to be able to use OLS
- Pick parameters of your model (b_0 , b_1 etc.), and calculate the *likelihood* of the data, given those parameters. We do this iteratively until we find the best parameters -- the ones that *maximize* the *likelihood* of your data.

GLM in R

```
glm(formula,  
    family = gaussian(link="",  
    data,  
    weights,  
    subset,  
    na.action,  
    start = NULL,  
    etastart,  
    mustart,  
    offset,  
    control = glm.control(..  
    model = TRUE,  
    method = "glm.fit",  
    x = FALSE,  
    y = TRUE,  
    contrasts = NULL, ...)
```

The **family** argument specifies the distribution. In R, families have default links.

Family	Default Link Function
binomial	(link = "logit")
gaussian	(link = "identity")
Gamma	(link = "inverse")
inverse.gaussian	(link = "1/mu^2")
poisson	(link = "log")
quasi	(link = "identity", variance = "constant")
quasibinomial	(link = "logit")
quasipoisson	(link = "log")

GLM in R

```
glm(y ~ X1+ X2 + X3 ,  
     family = binomial,  
     data = dataset)
```

Specify the model like
you would with `lm`

GLM in R

```
glm(y ~ X1+ X2 + X3 ,  
    family = binomial,  
    data = dataset)
```

Specify the distribution you're working with. When binary outcomes, we'll use the binomial.

GLM in R

```
glm(y ~ X1+ X2 + X3 ,  
     family = binomial,  
     data = dataset)
```

Specify your dataset.

- b_1 is the predicted change in the logit for a 1-unit change in X , holding the other predictors constant
- For a 1-unit change in X , holding other predictors constant, the odds that $Y = 1$ changes by e^{b_1}
 - e.g., $b_1 = .4$, $e^{.4} = 1.49$
- For fitted values, need to use entire equation $\hat{Y} = e^{b_0 + b_1 X_1}$
- Turn to probabilities by: $\frac{\text{odds}}{(1 + \text{odds})}$

Example

```
# 1 = not premature  
mortality
```

```
## # A tibble: 300 × 4  
##   Intelligence_Self Intelligence_Mate premature.d NOT.premature  
##           <dbl>           <dbl> <fct>           <dbl>  
## 1             22             19 normal             1  
## 2             22             18 normal             1  
## 3             21             21 normal             1  
## 4             22             17 normal             1  
## 5             19             18 normal             1  
## 6             19             20 premature            0  
## 7             16             18 normal             1  
## 8             15             11 premature            0  
## 9             16             21 normal             1  
## 10            19             22 normal             1  
## # i 290 more rows
```

```
death.1 <- lm(NOT.premature ~ Intelligence_Self , data = mort
summary(death.1)
```

```
##
## Call:
## lm(formula = NOT.premature ~ Intelligence_Self, data = mortality)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.9030   0.1084   0.1538   0.1907   0.3355
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.641769    0.098636   6.506 3.25e-10 ***
## Intelligence_Self 0.011357    0.005807   1.956  0.0514 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3745 on 298 degrees of freedom
## Multiple R-squared:  0.01267,    Adjusted R-squared:  0.009362
## F-statistic: 3.826 on 1 and 298 DF,  p-value: 0.05141
```

```
death.2 <- glm(NOT.premature ~ Intelligence_Self , data = mor
summary(death.2)
```

```
##
## Call:
## glm(formula = NOT.premature ~ Intelligence_Self, data = mortality)
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.641769   0.098636   6.506 3.25e-10 ***
## Intelligence_Self 0.011357   0.005807   1.956  0.0514 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.1402466)
##
##      Null deviance: 42.330  on 299  degrees of freedom
## Residual deviance: 41.793  on 298  degrees of freedom
## AIC: 266.05
##
## Number of Fisher Scoring iterations: 2
```

```
anova(death.1)
```

```
## Analysis of Variance Table
```

```
##
```

```
## Response: NOT.premature
```

```
##
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
--	----	--------	---------	---------	--------

## Intelligence_Self	1	0.537	0.53653	3.8256	0.05141
----------------------	---	-------	---------	--------	---------

## Residuals	298	41.793	0.14025		
--------------	-----	--------	---------	--	--

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
anova(death.2)
```

```
## Analysis of Deviance Table
```

```
##
```

```
## Model: gaussian, link: identity
```

```
##
```

```
## Response: NOT.premature
```

```
##
```

```
## Terms added sequentially (first to last)
```

```
##
```

```
##
```

```
##           Df Deviance Resid. Df Resid. Dev
```

```
## NULL                299      42.330
```

```
## Intelligence_Self  1  0.53653      298      41.793
```

```
death.3 <- glm(NOT.premature ~ Intelligence_Self,
               family = binomial, data = mortality)
summary(death.3)
```

```
##
## Call:
## glm(formula = NOT.premature ~ Intelligence_Self, family = binomial,
##      data = mortality)
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.28695    0.67490   0.425   0.6707
## Intelligence_Self 0.08012    0.04143   1.934   0.0532 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 273.53  on 299  degrees of freedom
## Residual deviance: 269.75  on 298  degrees of freedom
## AIC: 273.75
##
## Number of Fisher Scoring iterations: 4
```

Interpretation

For a 1-unit change in X, holding other predictors constant, the odds that Y = 1 changes by e^{b_1}

```
exp(1) ^ .08012
```

```
## [1] 1.083417
```

For every 1-unit increase in Intelligence, the odds of living increase by 8%

Interpreting Odds Ratios

What Does the OR Mean?

So, what does an OR mean? Here it is in plain language.

- An OR of 1.2 means there is a 20% increase in the odds of an outcome with a given exposure.
- An OR of 2 means there is a 100% increase in the odds of an outcome with a given exposure. Or this could be stated that there is a doubling of the *odds* of the outcome. Note, this is not the same as saying a doubling of the *risk*.
- An OR of 0.2 means there is an 80% decrease in the odds of an outcome with a given exposure.

Summary

- Odds Ratio is a measure of the strength of association with an exposure and an outcome.
 - $OR > 1$ means greater odds of association with the exposure and outcome.
 - $OR = 1$ means there is no association between exposure and outcome.
 - $OR < 1$ means there is a lower odds of association between the exposure and outcome.
- If the 95% confidence interval for the OR includes 1, the results are not statistically significant.
- OR and RR are not the same.
- OR always overestimate RR, but...
- OR approximates RR when the outcome is rare but markedly overestimates it as outcome exceeds 10%.

Specific Values?

What if you want the probability of being a premature death for a given level of Intelligence? (Now that we've run our model and have parameters...)

For fitted values, need to use entire equation

$$\hat{Y} = e^{b_0 + b_1 X_1}$$

```
# get odds with a given value of X (here 20)  
exp(1)^(0.28695 + (.08012*20))
```

```
## [1] 6.615067
```

```
# now get probability  
6.615067 / (1+6.615067)
```

Probit

We can have different link functions. When your response variable (DV) is truly binary -- the data generating process generates legit binary data -- logit is your pick.

What if your response variable is binary, but the underlying construct you are trying to measure is likely Gaussian? Ex: depressed vs. not depressed. But the underlying latent construct is continuous. More appropriate then is the **probit** link function.

Stack exchange thread if you're going down this route

Probit

```
death.4 <- glm(NOT.premature ~ Intelligence_Self,  
              family = binomial(link = "probit"), data = mortality)  
summary(death.4)
```

```
##  
## Call:  
## glm(formula = NOT.premature ~ Intelligence_Self, family = binomial(link =  
##      data = mortality))  
##  
## Coefficients:  
##              Estimate Std. Error z value Pr(>|z|)  
## (Intercept)      0.21961      0.38376   0.572   0.5671  
## Intelligence_Self 0.04513      0.02319   1.946   0.0516 .  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## (Dispersion parameter for binomial family taken to be 1)  
##  
##      Null deviance: 273.53  on 299  degrees of freedom  
## Residual deviance: 269.72  on 298  degrees of freedom  
## AIC: 273.72  
##
```

- You can include covariates
- You can have interactions...BUT

Interactions are super hard to interpret in logistic regression

- We have nonlinear mapping (S-function)
- Can express in terms of odds, probabilities, or logits
- Whether you observe an interaction *depends* on if you express the outcome in terms of odds, probabilities, or logits...You can get very different results!
- Ultimately, you might want to use other techniques

Weighted Least Squares

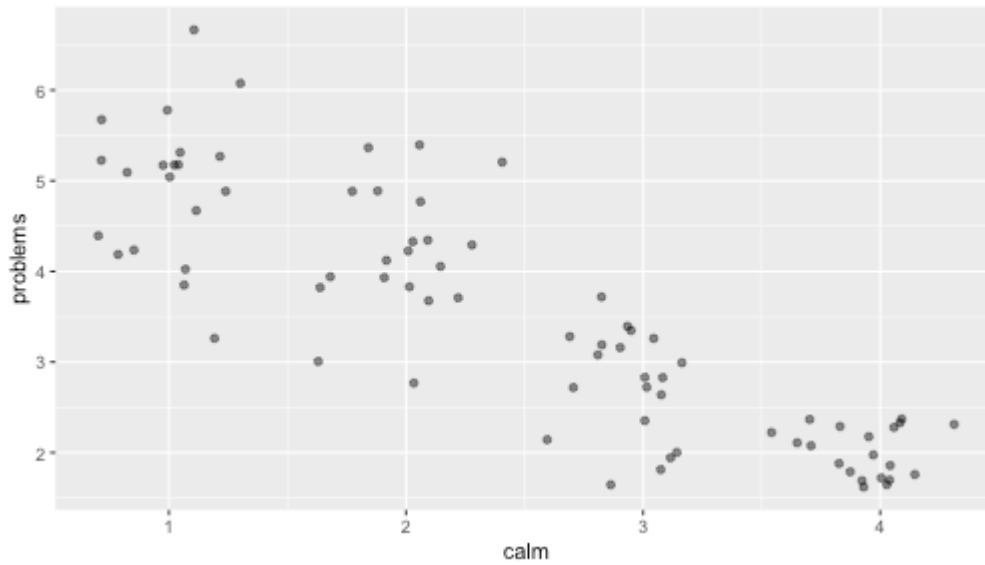
Estimation methods thus far

- OLS
- MLE

Why do we need another method?

Homoscedasticity

Homoscedasticity is the assumption that the variance of Y is constant across all levels of a predictor.



Do we meet this assumption?

Weighted least squares

Weighted least squares (WLS) is a commonly used remedial procedure for heteroscedasticity.

In an ordinary least squares (OLS) approach, each case in the dataset is given equal weight.

- WLS assigns each case a weight w_i , depending upon the precision of the observation of Y in that case.
- For observations for which the variance around the residuals around the regression line is low, the case is given a high weight.

OLS

Recall that an OLS estimation chooses values of b_0 and b_1 that minimizes the sum of squared residuals:

$$\min(\sum e_i^2) = \min \sum (Y_i - b_1 X_i - b_0)^2$$

WLS

Weights are taken into account, such that the values of b_0 and b_1 are chosen to minimize the sum of the **weighted** squared residuals:

$$\min(\sum w_i e_i^2) = \min \sum w_i (Y_i - b_1 X_i - b_0)^2$$

The value of the weights is the inverse of the conditional variance of the residuals corresponding to the specified value of X:

$$w_i = \frac{1}{\sigma_{Y-\hat{Y}|X}^2}$$

The value of $\sigma^2_{Y-\hat{Y}|X}$, the variance of the residuals in the population conditional on X , is not known and must be estimated.

A common procedure for estimating weights is to:

1. Estimate the usual OLS regression equation
2. Square the residuals
3. Regress the squared residuals onto X (x predicts squared residuals)

The weight is then estimated as the inverse of the predicted value for a case.

Our Data

Using our own data:

```
ols.model = lm(problems ~ calm, data = Data)
library(broom)
ols_aug = augment(ols.model)
head(ols_aug)
```

```
## # A tibble: 6 × 8
##   problems calm .fitted .resid   .hat .sigma .cooksd .std.resid
##   <dbl> <dbl>   <dbl> <dbl> <dbl> <dbl>   <dbl>      <dbl>
## 1         4 1.06     4.87 -0.872 0.0326 0.665 0.0296     -1.33
## 2         5 1.24     4.70  0.304 0.0278 0.671 0.00306      0.462
## 3         6 1.30     4.63  1.37  0.0263 0.653 0.0581      2.07
## 4         4 0.852     5.08 -1.08  0.0391 0.660 0.0558     -1.66
## 5         5 1.12     4.82  0.180 0.0311 0.672 0.00120      0.273
## 6         5 0.715     5.22 -0.223 0.0438 0.672 0.00266     -0.341
```

```
# square residuals
ols_aug$resid_sq = ols_aug$.resid^2
# regress squared resid on predictor
weight.mod = lm(resid_sq ~ calm, data = ols_aug)
```

```
coef(ols.model)
```

```
## (Intercept)      calm
##    5.941940   -1.005699
```

```
coef(weight.mod)
```

```
## (Intercept)      calm
##    1.0739505  -0.2591579
```

```
# extract predicted values
pred.resid = predict(weight.mod)
head(pred.resid)
```

```
##           1           2           3           4           5           6
## 0.7982546 0.7527584 0.7366885 0.8530291 0.7849316 0.8886457
```

```
# find inverse of predicted values
# use absolute value if some of your predicted values are neg
est.weights = 1/abs(pred.resid)
```


```
wls.model = lm(problems ~ calm, data = Data, weights = est.we
```

```
tidy(ols.model)
```

```
## # A tibble: 2 × 5
##   term          estimate std.error statistic  p.value
##   <chr>          <dbl>     <dbl>     <dbl>    <dbl>
## 1 (Intercept)    5.94      0.182      32.6 3.54e-47
## 2 calm         -1.01      0.0675     -14.9 1.58e-24
```

```
tidy(wls.model)
```

```
## # A tibble: 2 × 5
##   term          estimate std.error statistic  p.value
##   <chr>          <dbl>     <dbl>     <dbl>    <dbl>
## 1 (Intercept)    5.64      0.182      30.9 1.60e-45
## 2 calm         -0.884     0.0451     -19.6 7.58e-32
```



When to use WLS

WLS is a **robust** method of estimation. **Robust statistics** are statistics with good performance for data drawn from a wide range of probability distributions, especially for distributions that are not normal." (*Wikipedia*)

Use when...

- Dealing with heteroscedasticity
- You know already that points should not be treated equally (some should be weighed more than others); do you know that some of your data was measured with less error?

Totally Random + HW Assignment

Checking Assumptions

```
library(performance)
library(palmerpenguins)

model1 = lm(body_mass_g ~ bill_length_mm, data = penguins)
summary(model1)
```

```
##
## Call:
## lm(formula = body_mass_g ~ bill_length_mm, data = penguins)
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-1762.08	-446.98	32.59	462.31	1636.86

```
##
## Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	362.307	283.345	1.279	0.202
bill_length_mm	87.415	6.402	13.654	<2e-16 ***


```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 645.4 on 340 degrees of freedom
```

Checking Assumptions

```
check_model(x = model1)
```


Posterior Predictive Check

Model-predicted lines should resemble observed data




Linearity

Reference line should be flat and horizontal




Homogeneity of Variance

Reference line should be flat and horizontal




Influential Observations

Points should be inside the contour lines



Normality of Residuals

Points should fall along the line



End of semester

- Tuesday = Resampling methods, namely bootstrapping
- Thursday = Machine learning; READ THE YARKONI & WESTFALL PAPER!
- Tuesday = Wrapping up, future directions, review
- Thursday = Exam 3. Oral exams will open after that.